



TEXAS TECH UNIVERSITY™




Lecture 8: Probing Moderation

Kyle M. Lang

*Institute for Measurement, Methodology, Analysis & Policy
Texas Tech University
Lubbock, TX*

2016 Stats Camp




- Probing moderation via centering
- Alternative probing strategies
- Confidence bands for simple slopes

Starting Point



Last time, we fit the following model:



We probed the interaction with the *Pick-a-point* approach.

- Choose interesting values of the moderator Z
- Check the significance of the focal effect $X \rightarrow Y$ at the values we choose for Z .
- Gives us an idea of where in Z 's distribution the focal effect is/is not significant.

Previously, we manually calculated the all of the quantities we needed, including a SE for the conditional focal effect.

- There is a simpler way: CENTERING

Centering transforms a variable by subtracting a constant (e.g., the variable's mean) from each observation of the variable

- The most familiar form of center is *mean centering*
- We can center on any value
 - When probing interactions, we can center Z on the interesting values we choose during the pick-a-point approach
 - Running the model with Z centered on specific values automatically provides tests of the simple slope conditional on those values of Z

Say we want to do a simple slopes analysis to test the conditional effect of X on Y at three levels of $Z = \{Z_1, Z_2, Z_3\}$.

Then, all we need to do is fit the following three models:

$$Y = \alpha + \beta_1 X + \beta_2(Z - Z_1) + \beta_3 X(Z - Z_1) + e$$

$$Y = \alpha + \beta_1 X + \beta_2(Z - Z_2) + \beta_3 X(Z - Z_2) + e$$

$$Y = \alpha + \beta_1 X + \beta_2(Z - Z_3) + \beta_3 X(Z - Z_3) + e$$

The default output for β_1 provides tests of the simple slopes.

Example



```
## Read in the data:
dataDir <- "../data/"
fileName <- "nlsyData.rds"
dat1 <- readRDS(paste0(dataDir, fileName))
## Moderated Model:
out1 <- lm(depress1 ~ ratio1*perception1, data = dat1)
summary(out1)
```

Call:

```
lm(formula = depress1 ~ ratio1 * perception1, data = dat1)
```

Residuals:

Min	1Q	Median	3Q	Max
-2.1218	-0.2830	0.0881	0.3515	1.0835

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	2.41374	0.10375	23.265	< 2e-16	***
ratio1	0.38126	0.04571	8.341	< 2e-16	***
perception1	0.11083	0.02941	3.768	0.000165	***
ratio1:perception1	-0.07702	0.01214	-6.343	2.37e-10	***

Example



```
Signif. codes:  0      ***      0.001      **      0.01      *      0.05
                .      0.1          1
```

```
Residual standard error: 0.5086 on 8980 degrees of freedom
Multiple  $R^2$ : 0.01426, Adjusted  $R^2$ : 0.01393
F-statistic: 43.3 on 3 and 8980 DF, p-value: < 2.2e-16
```

```
## Compute critical values of Z:
zMean <- mean(dat1$perception1)
zSD <- sd(dat1$perception1)
dat1$zCen <- dat1$perception1 - zMean
dat1$zHigh <- dat1$perception1 - (zMean + zSD)
dat1$zLow <- dat1$perception1 - (zMean - zSD)
## Test simple slopes:
out2.1 <- lm(depress1 ~ ratio1*zLow, data = dat1)
summary(out2.1)
```


Example



Call:

```
lm(formula = depress1 ~ ratio1 * zLow, data = dat1)
```

Residuals:

Min	1Q	Median	3Q	Max
-2.1218	-0.2830	0.0881	0.3515	1.0835

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	2.68561	0.04066	66.052	< 2e-16	***
ratio1	0.19234	0.01888	10.190	< 2e-16	***
zLow	0.11083	0.02941	3.768	0.000165	***
ratio1:zLow	-0.07702	0.01214	-6.343	2.37e-10	***

Signif. codes:	0	***	0.001	**	0.01	*	0.05
	.	0.1	1				

Residual standard error: 0.5086 on 8980 degrees of freedom

Multiple R^2 : 0.01426, Adjusted R^2 : 0.01393

F-statistic: 43.3 on 3 and 8980 DF, p-value: < 2.2e-16

Example



```
out2.2 <- lm(depress1 ~ ratio1*zCen, data = dat1)
summary(out2.2)
```

Call:

```
lm(formula = depress1 ~ ratio1 * zCen, data = dat1)
```

Residuals:

Min	1Q	Median	3Q	Max
-2.1218	-0.2830	0.0881	0.3515	1.0835

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	2.76845	0.03076	89.998	< 2e-16 ***
ratio1	0.13477	0.01345	10.021	< 2e-16 ***
zCen	0.11083	0.02941	3.768	0.000165 ***
ratio1:zCen	-0.07702	0.01214	-6.343	2.37e-10 ***

Signif. codes:	0	***	0.001	**	0.01	*	0.05
	.	0.1	1				

Residual standard error: 0.5086 on 8980 degrees of freedom
Multiple R^2 : 0.01426, Adjusted R^2 : 0.01393

Example



```
F-statistic: 43.3 on 3 and 8980 DF, p-value: < 2.2e-16
```

```
out2.3 <- lm(depress1 ~ ratio1*zHigh, data = dat1)
summary(out2.3)
```

Call:

```
lm(formula = depress1 ~ ratio1 * zHigh, data = dat1)
```

Residuals:

Min	1Q	Median	3Q	Max
-2.1218	-0.2830	0.0881	0.3515	1.0835

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	2.85128	0.03472	82.115	< 2e-16	***
ratio1	0.07721	0.01305	5.919	3.36e-09	***
zHigh	0.11083	0.02941	3.768	0.000165	***
ratio1:zHigh	-0.07702	0.01214	-6.343	2.37e-10	***

Signif. codes:	0	***	0.001	**	0.01	*	0.05
	.	0.1	1				

Example



Residual standard error: 0.5086 on 8980 degrees of freedom
Multiple R^2 : 0.01426, Adjusted R^2 : 0.01393
F-statistic: 43.3 on 3 and 8980 DF, p-value: $< 2.2e-16$

Compare Approaches



The manual and the centering approaches give identical answers, barring rounding errors with the manual approach:

	Z Low	Z Center	Z High
Manual	0.192335	0.134774	0.077213
Centering	0.192335	0.134774	0.077213

Table: Simple Slopes

	Z Low	Z Center	Z High
Manual	0.018875	0.013449	0.013045
Centering	0.018875	0.013449	0.013045

Table: Standard Errors

A Few Comments on Centering



You will often hear mean centering touted as absolutely necessary or absolutely unnecessary for moderation analysis.

Both sides are partially correct.

Two effects are usually ascribed to mean centering in moderation analysis:

1. Improved interpretation of the conditional effects
2. Reduced multicollinearity between lower-order effects and the interaction term

A Few Comments on Centering



Mean center absolutely *does* have the potential to improve parameter interpretations

- When $X = 0$ or $Z = 0$ are not sensible values, centering is necessary for any plausible interpretation of β_1 or β_2 .

Mean centering *can* remove collinearity between lower-order terms and the interaction term

- **BUT**, we don't care

We can get a better sense of what's going on with a synthetic example.

Example



```
n ← 10000  
x ← rnorm(n, 10, 1)  
z ← rnorm(n, 20, 2)  
xz ← x*z  
cor(x, xz)
```


```
[1] 0.7065279
```

```
xc ← x - mean(x)  
zc ← z - mean(z)  
xzc ← xc*zc  
cor(xc, xzc)
```


```
[1] -0.002515431
```


Example

Uncentered



Centered



Example



```
y ← 5*x + 5*z + 2*xz + rnorm(n, 0, 0.5)
out3.1 ← lm(y ~ x*z)
out3.2 ← lm(y ~ xc*z)
out3.3 ← lm(y ~ xc*zc)
summary(out3.1)$coef
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	0.6385057	0.515435755	1.238769	0.2154603
x	4.9437053	0.051277645	96.410537	0.0000000
z	4.9697606	0.025657191	193.698546	0.0000000
x:z	2.0026657	0.002552174	784.690103	0.0000000

```
summary(out3.2)$coef
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	50.081375	0.050828516	985.30077	0
xc	4.943705	0.051277645	96.41054	0
z	24.998774	0.002529088	9884.50339	0
xc:z	2.002666	0.002552174	784.69010	0

```
summary(out3.3)$coef
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	550.015671	0.005045344	109014.5048	0
xc	44.993721	0.005069554	8875.2812	0
zc	24.998774	0.002529088	9884.5034	0
xc:zc	2.002666	0.002552174	784.6901	0

Example



```
sum(out3.1$fitted - out3.3$fitted)
```

```
[1] -2.785328e-12
```

```
summary(out3.1)$r.squared
```

```
[1] 0.9999439
```

```
summary(out3.3)$r.squared
```

```
[1] 0.9999439
```

Example



```
summary(lm(y ~ x*zLow))$coef
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	90.110684	0.072088368	1250.0031	0
x	40.998333	0.007175983	5713.2703	0
zLow	4.969761	0.025657191	193.6985	0
x:zLow	2.002666	0.002552174	784.6901	0

```
summary(lm(y ~ xc*zLow))$coef
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	500.142247	0.007135321	70093.8655	0
xc	40.998333	0.007175983	5713.2703	0
zLow	24.998774	0.002529088	9884.5034	0
xc:zLow	2.002666	0.002552174	784.6901	0

Example



```
summary(lm(y ~ x*zCen))$coef
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	100.025529	0.050952002	1963.1325	0
x	44.993721	0.005069554	8875.2812	0
zCen	4.969761	0.025657191	193.6985	0
x:zCen	2.002666	0.002552174	784.6901	0

```
summary(lm(y ~ xc*zCen))$coef
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	550.015671	0.005045344	109014.5048	0
xc	44.993721	0.005069554	8875.2812	0
zCen	24.998774	0.002529088	9884.5034	0
xc:zCen	2.002666	0.002552174	784.6901	0

Example



```
summary(lm(y ~ x*zHigh))$coef
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	109.940374	0.072358107	1519.3926	0
x	48.989108	0.007194189	6809.5387	0
zHigh	4.969761	0.025657191	193.6985	0
x:zHigh	2.002666	0.002552174	784.6901	0

```
summary(lm(y ~ xc*zHigh))$coef
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	599.889094	0.007135453	84071.6162	0
xc	48.989108	0.007194189	6809.5387	0
zHigh	24.998774	0.002529088	9884.5034	0
xc:zHigh	2.002666	0.002552174	784.6901	0

QUESTION: Okay, so what about our example analysis? Should we center the predictors in this model:

$$Depress = \alpha + \beta_1 Ratio + \beta_2 Perception + \beta_3 Ratio \times Perception + e$$

QUESTION: Okay, so what about our example analysis? Should we center the predictors in this model:

$$Depress = \alpha + \beta_1 Ratio + \beta_2 Perception + \beta_3 Ratio \times Perception + e$$

ANSWER: Yes.

QUESTION: Okay, so what about our example analysis? Should we center the predictors in this model:

$$Depress = \alpha + \beta_1 Ratio + \beta_2 Perception + \beta_3 Ratio \times Perception + e$$

ANSWER: Yes.

QUESTION MARK II: Why?

QUESTION: Okay, so what about our example analysis? Should we center the predictors in this model:

$$Depress = \alpha + \beta_1 Ratio + \beta_2 Perception + \beta_3 Ratio \times Perception + e$$

ANSWER: Yes.

QUESTION MARK II: Why?

ANSWER THE SECOND: Because a *Weight:Height* ratio of zero is nonsensical and zero is outside the range of *Perception*.

Example



```
dat1$ratioC ← dat1$ratio1 - mean(dat1$ratio1)
dat1$perceptionC ← dat1$perception1 - mean(dat1$perception1)
## Moderated Model:
out1 ← lm(depress1 ~ ratioC*perceptionC, data = dat1)
summary(out1)$coef
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	3.07996908	0.005830947	528.210782	0.000000e+00
ratioC	0.13477427	0.013448801	10.021285	1.632952e-23
perceptionC	-0.06719252	0.008311533	-8.084252	7.065839e-16
ratioC:perceptionC	-0.07701956	0.012142973	-6.342726	2.366288e-10

```
## Compute critical values of Z:
zMean ← mean(dat1$perceptionC)
zSD ← sd(dat1$perceptionC)
dat1$zCen ← dat1$perceptionC - zMean
dat1$zHigh ← dat1$perceptionC - (zMean + zSD)
dat1$zLow ← dat1$perceptionC - (zMean - zSD)
```

Example



```
## Test simple slopes:
```

```
out2.1 <- lm(depress1 ~ ratioC*zLow, data = dat1)
summary(out2.1)$coef
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	3.13018574	0.008490579	368.665759	0.000000e+00
ratioC	0.19233521	0.018875273	10.189798	2.983313e-24
zLow	-0.06719252	0.008311533	-8.084252	7.065839e-16
ratioC:zLow	-0.07701956	0.012142973	-6.342726	2.366288e-10

```
out2.2 <- lm(depress1 ~ ratioC*zCen, data = dat1)
summary(out2.2)$coef
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	3.07996908	0.005830947	528.210782	0.000000e+00
ratioC	0.13477427	0.013448801	10.021285	1.632952e-23
zCen	-0.06719252	0.008311533	-8.084252	7.065839e-16
ratioC:zCen	-0.07701956	0.012142973	-6.342726	2.366288e-10

```
out2.3 <- lm(depress1 ~ ratioC*zHigh, data = dat1)
summary(out2.3)$coef
```

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	3.02975242	0.008548655	354.412740	0.000000e+00
ratioC	0.07721333	0.013045300	5.918862	3.360611e-09
zHigh	-0.06719252	0.008311533	-8.084252	7.065839e-16
ratioC:zHigh	-0.07701956	0.012142973	-6.342726	2.366288e-10

The pick-a-point approach is nice due to its simplicity and ease of interpretation, but the points we choose are totally arbitrary.

- We may be missing important nuances that occur in some of the areas of Z 's distribution that we *did not* pick.

The pick-a-point approach is nice due to its simplicity and ease of interpretation, but the points we choose are totally arbitrary.

- We may be missing important nuances that occur in some of the areas of Z 's distribution that we *did not* pick.

The *Johnson-Neyman* technique is an alternative approach that removes the arbitrary choices necessary for pick-a-point.

- Johnson-Neyman finds the *region of significance* wherein the conditional effect of X on Y is statistically significant
- Inverts the pick-a-point approach to find what cut-points on the moderator correspond to a critical t value for the conditional β_1 .

With pick-a-point, we:

1. Choose conditional values of Z , say Z_1
2. Calculate the simple slope SS_1 and standard error SE_{SS_1} associated with Z_1
3. Test SS_1 for significance via a simple Wald-type test:

$$t = \frac{SS_1}{SE_{SS_1}} \quad (1)$$

With Johnson-Neyman, we:

1. Choose an α level for our test and the corresponding critical value of t , say $t_{crit} = 1.96$ to give $\alpha = 0.05$ in large samples.
2. Re-arrange Equation 1 into the following quadratic form:

$$t_{crit}^2 SE_{SS}^2 - SS^2 = 0 \quad (2)$$

3. Solve Equation 2 to find the two values of Z that produce critical t statistics for the conditional focal effect.

Johnson-Neyman Technique



The roots produced by the Johnson-Neyman technique delineate the *region of significance*.

- The conditional effect of X on Y is either significant everywhere between these two points or everywhere outside of these two points.
- If only one of the points falls within the observed range of Z , ignore the other point
 - The region of significant is either everywhere above or below the legal root
- If neither of the roots fall within the observed range of Z then, either:
 1. The focal effect is significant across the entire range of Z , or
 2. The focal effect is not significant anywhere within the range of Z

Recall the formula for a simple slope:

$$SS = \beta_1 + \beta_3 Z$$

From a graphical perspective, we can think about SS in, at least, two different ways:

1. As a weight for X that we can use to get plots of the conditional effect of X on Y at different levels of Z .
2. We can also consider how SS , itself, smoothly changes as a function of Z .

The latter perspective embodies the spirit of the Johnson-Neyman technique.

A natural quantity to consider is a confidence interval for SS :

$$CI_{SS} = SS \pm t_{crit} \cdot SE_{SS}$$

Last time, we computed a few such intervals for the interesting values of Z we chose for the pick-a-point analysis.

When doing Johnson-Neyman, we can considering the values of CI_{SS} for the entire range of Z .

- These CI values define the *confidence bands* of SS and show, for any value of Z , the corresponding CI for SS
 - As a result, we can immediately check any value of Z for a significant simple slope

Example




Implementing the Johnson-Neyman technique by hand is a pain, but we can easily do so by using the **rockchalk** package in R

```
par(family = "serif", cex = 0.75)
library(rockchalk)
## First we need to create a 'plotSlopes' object:
plotOut <- plotSlopes(model = out1,
                      plotx = "ratioC",
                      modx = "perceptionC",
                      plotPoints = FALSE)
## Then we modify 'plotOut' to get the J-N test:
testOut <- testSlopes(plotOut)
```

Values of perceptionC OUTSIDE this interval:

```
      lo      hi
1.327479 2.452667
cause the slope of (b1 + b2*perceptionC)ratioC to be
statistically significant
```

Example



We can see the significance boundaries by extracting the roots from 'testOut'

```
testOut$jn$roots
```

```
      lo      hi  
1.327479 2.452667
```

We can plot the result:

```
par(cex = 0.75, family = "serif")  
plot(testOut)
```

Example

