

Missing Data in SEM

Introduction to SEM with Lavaan



**Utrecht
University**

Kyle M. Lang

Department of Methodology & Statistics
Utrecht University

Outline

Missing Data Mechanisms



What are Missing Data?

Missing data are empty cells in a dataset where there should be observed values.

- The missing cells correspond to true population values, but we haven't observed those values.



What are Missing Data?

Missing data are empty cells in a dataset where there should be observed values.

- The missing cells correspond to true population values, but we haven't observed those values.

Not every empty cell is a missing datum.

- Quality-of-life ratings for dead patients in a mortality study
- Firm profitability after the company goes out of business
- Self-reported severity of menstrual cramping for men
- Empty blocks of data following “gateway” items



A Little Notation

$Y :=$ An $N \times P$ Matrix of Arbitrary Data

$Y_{mis} :=$ The *missing* part of Y

$Y_{obs} :=$ The *observed* part of Y

$R :=$ An $N \times P$ response matrix

$M :=$ An $N \times P$ missingness matrix

The R and M matrices are complementary.

- $r_{np} = 1$ means y_{np} is observed; $m_{np} = 1$ means y_{np} is missing.
- $r_{np} = 0$ means y_{np} is missing; $m_{np} = 0$ means y_{np} is observed.
- M_p is the *missingness* of Y_p .

Example

```
## Load some useful packages:
```

```
library(dplyr)
```

```
library(naniar)
```

```
library(ggmice)
```

```
## Read in some data:
```

```
bfi <- readRDS("../data/bfi_datasets.rds")$incomplete %>%  
  select(-matches("N\\d|C\\d|E\\d|male"))
```

```
## Compute the variablewise proportions of missing data:
```

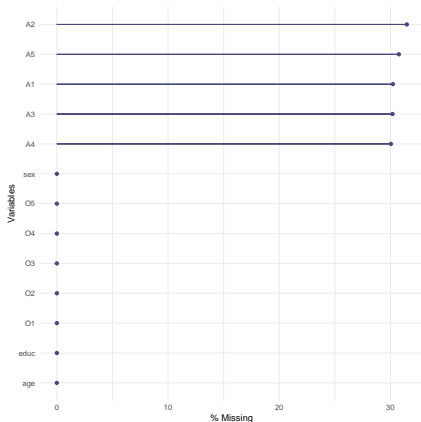
```
bfi %>% is.na() %>% colMeans() %>% round(2)
```

A1	A2	A3	A4	A5	01	02	03	04	05	age	sex	educ
0.30	0.31	0.30	0.30	0.31	0.00	0.00	0.00	0.00	0.00	0.00	0.00	0.00

Example

Visualize the percentages missing via **naniar::gg_miss_var()**.


```
gg_miss_var(bfi, show_pct = TRUE)
```



Example

Visualize the missing data patterns via **ggmice::plot_pattern()**.

```
plot_pattern(bfi)
```



Example

In **lavaan**, we can directly fit a model with incomplete data.

```
library(lavaan)

## Specify the measurement model:
mod <- "
agree =~ A1 + A2 + A3 + A4 + A5
open  =~ O1 + O2 + O3 + O4 + O5
"

## Estimate the model:
out <- cfa(mod, data = bfi, std.lv = TRUE)
```

Example

The model will estimate, just fine...

lavaan 0.6-11 ended normally after 21 iterations

Latent Variables:

	Estimate	Std.Err	z-value	P(> z)
agree =~				
A1	0.561	0.071	7.866	0.000
A2	-0.725	0.056	-13.055	0.000
A3	-0.901	0.063	-14.383	0.000
A4	-0.599	0.074	-8.103	0.000
A5	-0.781	0.059	-13.264	0.000
open =~				
O1	0.490	0.055	8.983	0.000
O2	-0.844	0.079	-10.649	0.000
O3	0.786	0.058	13.574	0.000
O4	0.311	0.055	5.699	0.000
O5	-0.874	0.073	-11.987	0.000

Example

Covariances:

	Estimate	Std.Err	z-value	P(> z)
agree ~~				
open	-0.254	0.060	-4.204	0.000

Variances:

	Estimate	Std.Err	z-value	P(> z)
.A1	1.662	0.114	14.593	0.000
.A2	0.781	0.067	11.650	0.000
.A3	0.865	0.086	10.050	0.000
.A4	1.776	0.122	14.517	0.000
.A5	0.863	0.075	11.430	0.000
.01	0.880	0.064	13.823	0.000
.02	1.704	0.133	12.776	0.000
.03	0.658	0.072	9.150	0.000
.04	0.976	0.065	15.020	0.000
.05	1.288	0.112	11.460	0.000

Example

But not everything is as it seems.

Estimator	ML	
Optimization method	NLMINB	
Number of model parameters	21	
	Used	Total
Number of observations	492	2800

Model Test User Model:

Test statistic	79.928
Degrees of freedom	34
P-value (Chi-square)	0.000

Default Approach

Like most software packages, **lavaan** will default to *complete case analysis* when asked to analyze incomplete data.

- In the absence of user input, this is a sensible option.
- That doesn't mean you should actually use deletion to treat the missing data in your analysis.



Default Approach

Like most software packages, **lavaan** will default to *complete case analysis* when asked to analyze incomplete data.

- In the absence of user input, this is a sensible option.
- That doesn't mean you should actually use deletion to treat the missing data in your analysis.

Complete case analysis has two major problems.

1. Throws out useful information (potentially a lot of information)
2. Probably biases parameter estimates.

To understand the second point, we need to discuss *missing data mechanisms*.

MISSING DATA MECHANISMS



Missing Data Mechanisms

Missing Completely at Random (MCAR)

- $P(R|Y_{mis}, Y_{obs}) = P(R)$
- Missingness is unrelated to any study variables.

Missing at Random (MAR)

- $P(R|Y_{mis}, Y_{obs}) = P(R|Y_{obs})$
- Missingness is related to only the *observed* parts of study variables.

Missing not at Random (MNAR)

- $P(R|Y_{mis}, Y_{obs}) \neq P(R|Y_{obs})$
- Missingness is related to the *unobserved* parts of study variables.



Simulate Some Toy Data

```
nObs <- 5000 # Sample Size
pm    <- 0.3  # Proportion Missing

sigma <- matrix(c(1.0, 0.5, 0.3,
                  0.5, 1.0, 0.0,
                  0.3, 0.0, 1.0),
                ncol = 3)
tmp <- rmvnorm(nObs, c(0, 0, 0), sigma)

x0 <- tmp[, 1]
y0 <- tmp[, 2]
z0 <- tmp[, 3]

cor(y0, x0) # Check correlation between X and Y

[1] 0.4997145
```

MCAR Example


```
## Simulate MCAR Missingness:
mVec <- sample(1 : length(y0), size = pm * length(y0))

yMcar      <- y0
yMcar[mVec] <- NA

cor(yMcar, x0, use = "pairwise") # Look at correlation

[1] 0.5195767
```

MCAR Example



MAR Example

```
## Simulate MAR Missingness:
mVec <- x0 < quantile(x0, probs = pm)
mean(mVec)

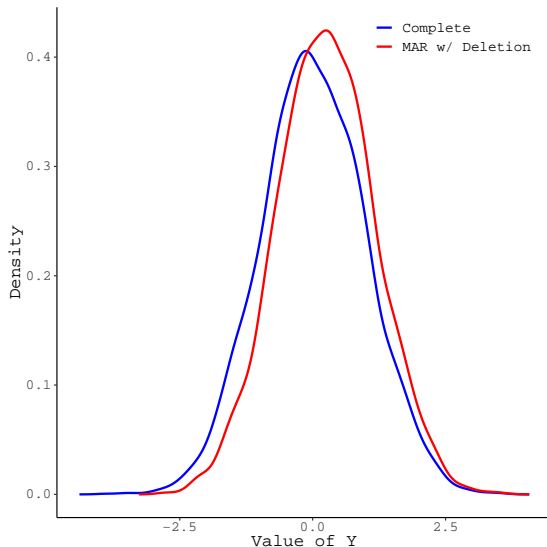
[1] 0.3

yMar      <- y0
yMar[mVec] <- NA

cor(yMar, x0, use = "pairwise") # Not looking so good :(

[1] 0.3822143
```

MAR Example



MNAR Example

```
## Simulate MNAR Missingness:
mVec <- y0 < quantile(y0, probs = pm)
mean(mVec)

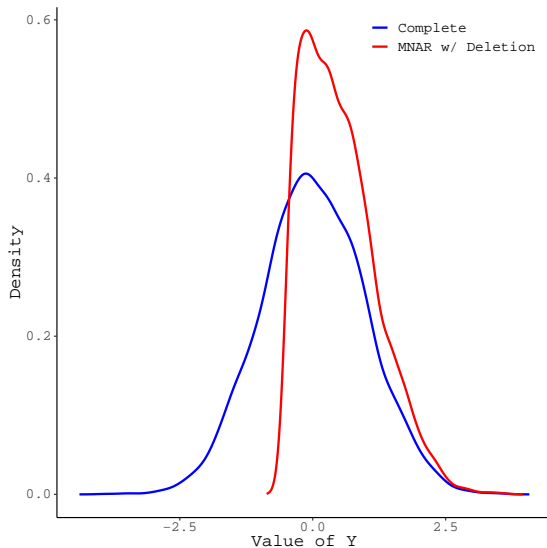
[1] 0.3

yMnar <- y0
yMnar[mVec] <- NA

cor(yMnar, x0, use = "pairwise") # Hmm...looks pretty bad.

[1] 0.3902962
```

MNAR Example



Effects of Deletion

As we saw in the preceding plots, excluding incomplete cases usually alters the variables' distributions.


- The statistics upon which we base our analyses generally summarize these distributions.
- Problems with the distributions show up as bias in the results of our analyses.

```
diabetes1 <- diabetes2 <- readRDS("../data/diabetes.rds")

mVec <- simLogisticMissingness0(data      = diabetes1,
                                pm        = 0.3,
                                preds     = c("bmi", "bp"),
                                stdData  = TRUE)$r

diabetes2[mVec, "glu"] <- NA
```


Example



Example

```
diabetes1 %>% select(bmi, glu, bp) %>% cor()
```

	bmi	glu	bp
bmi	1.0000000	0.38868	0.3954109
glu	0.3886800	1.00000	0.3904300
bp	0.3954109	0.39043	1.0000000

```
diabetes2 %>% select(bmi, glu, bp) %>% cor(use = "complete")
```

	bmi	glu	bp
bmi	1.0000000	0.2566595	0.2052338
glu	0.2566595	1.0000000	0.3011547
bp	0.2052338	0.3011547	1.0000000

Example

```
mean(diabetes1$glu)
```

```
[1] 91.26018
```

```
mean(diabetes2$glu, na.rm = TRUE)
```

```
[1] 88.86424
```

```
var(diabetes1$glu)
```

```
[1] 132.1657
```

```
var(diabetes2$glu, na.rm = TRUE)
```

```
[1] 115.8254
```

Good Methods

Multiple Imputation (MI)


- Replace the missing values with M plausible estimates
 - Essentially, a repeated application of stochastic regression imputation (with a particular type of regression model)
 - Produces unbiased parameter estimates and predictions
 - Produces “correct” standard errors, CIs, and prediction intervals
 - Very, very flexible
 - Computationally expensive



Good Methods

Full Information Maximum Likelihood (FIML)

- Adjust the objective function to only consider the observed parts of the data
 - Models are directly estimated in the presence of missing data
 - The predictors of nonresponse must be included in the model, somehow
 - Unless you write your own optimization program, FIML is only available for certain types of models
 - In linear regression models, FIML cannot treat missing data on predictors (if the predictors are taken as fixed)



Good Methods

What happens when we apply MI to our previous MAR example?

```
## Estimate imputation model:
miceOut <- mice(data      = data.frame(y = yMar, x = x0),
               m          = 25,
               maxit      = 1,
               method     = "norm",
               printFlag  = FALSE)


## Estimate and pool M correlations:
with(miceOut, cor(y, x))$analyses %>% unlist() %>% mean()

[1] 0.504661
```

The MI-based parameter estimate looks good.

- MI produces unbiased estimates of the parameter when data are MAR.

Good Methods



Good Methods

What about applying MI to our MNAR example?

```
## Estimate imputation model:
miceOut <- mice(data      = data.frame(y = yMnar, x = x0),
               m          = 25,
               maxit      = 1,
               method     = "norm",
               printFlag  = FALSE)

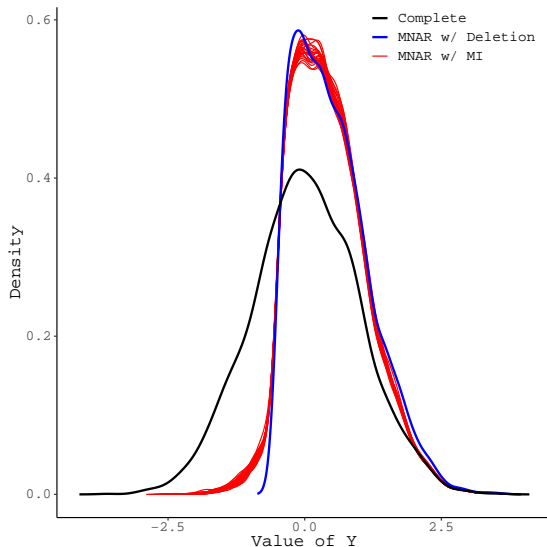
## Estimate and pool M correlations:
with(miceOut, cor(y, x))$analyses %>% unlist() %>% mean()

[1] 0.4116519
```

The MI-based parameter estimate is still biased.

- MI cannot correct bias in parameter estimates when data are MNAR.

Good Methods



MI Example

```
## The mice package does MI:  
library(mice)  
  
## Multiply impute the missing data:  
miceOut <- mice(data      = diabetes2,  
                 m        = 25,  
                 maxit     = 1,  
                 printFlag = FALSE,  
                 method    = "norm")
```



MI Example

```
## Complete data:
```

```
diabetes1 %>% select(bmi, glu, bp) %>% cor()
```

	bmi	glu	bp
bmi	1.0000000	0.38868	0.3954109
glu	0.3886800	1.00000	0.3904300
bp	0.3954109	0.39043	1.0000000

```
## MI:
```

```
pooledCorMat(miceOut, c("bmi", "glu", "bp"))
```

	bmi	glu	bp
bmi	1.0000000	0.3135162	0.3954109
glu	0.3135162	1.0000000	0.3563903
bp	0.3954109	0.3563903	1.0000000

MI Example

```
mean(diabetes1$glu)
```

```
[1] 91.26018
```

```
with(miceOut, mean(glu))$analyses %>% unlist() %>% mean()
```

```
[1] 90.61747
```

```
var(diabetes1$glu)
```

```
[1] 132.1657
```

```
with(miceOut, var(glu))$analyses %>% unlist() %>% mean()
```

```
[1] 123.3748
```

FIML Example

```
fit <- diabetes2 %>%  
  select(bmi, glu, bp) %>%  
  lavCor(missing = "fiml", output = "sampstat")  
  
## Complete data:  
diabetes1 %>% summarize(mean = mean(glu), var = var(glu))  
  
      mean      var  
1 91.26018 132.1657  
  
## FIML:  
fit %$% c(mean = mean[["glu"]], var = cov[["glu", "glu"]])  
  
      mean      var  
90.82487 125.27146
```

FIML Example

```
diabetes1 %>% select(bmi, glu, bp) %>% cor() %>% round(3)
```

```
      bmi    glu    bp  
bmi 1.000 0.389 0.395  
glu 0.389 1.000 0.390  
bp  0.395 0.390 1.000
```

```
fit$cov %>% cov2cor()
```

```
      bmi    glu    bp  
bmi 1.000  
glu 0.357 1.000  
bp  0.395 0.386 1.000
```

```
diabetes2 %>% select(bmi, glu, bp) %>% cor(use = "complete") %>% round(3)
```

```
      bmi    glu    bp  
bmi 1.000 0.257 0.205  
glu 0.257 1.000 0.301  
bp  0.205 0.301 1.000
```

References
